

Deep Learning 101

Astro Hack Week 2019 - Cambridge UK

Gabriella Contardo (Flatiron Institute, Center for Computational Astrophysics)
gcontardo@flatironinstitute.org

Francois Lannuse (UC Berkeley) flannuse@berkeley.edu

Outline of tutorial part I

- Machine learning principles
- *Why deep learning?*
- Neural networks:
 - Components, training
 - Fully Connected / Multi-Layer Perceptron (MLP)
 - Convolutional neural networks
- Build your own model in tensorflow !

A bit of definitions

Machine Learning (ML) :

- Relying on data analysis to automate model building
- Model learns from data, extracts patterns, and makes some 'prediction'
- E.g. SVM, Random Forest, KNN, Deep learning....

Deep Learning :

- Subgroup of ML based on artificial neural networks
- Strongly rely on “pattern” extraction / representation learning

Learning can be **supervised**, **unsupervised**, or **semi-supervised**.

Some Machine Learning principles

3 main “components”:

- Data
- Machine or model
- Criterion (learning and evaluation)

Goal : Automatically extract relevant information from the data that *generalize* well to infer on new data (‘test’ or evaluation data from similar distribution, but different from training examples)

Learn a model on training data to optimize some criterion, then infer on test data.

Supervised Learning

Training data is composed of

- Examples (inputs)
- 'Targets' --or label(s)-- (i.e outputs) for each example.

E.g. image classification.



Different type of targets => different learning problem:

- Limited set of values, categorical : classification (multi-class, multi-label,...)
- Continuous values : regression
- Other type e.g. noisy targets, weak supervision etc.

Goal: Find an approximation of the labeling process.

What are neural networks?

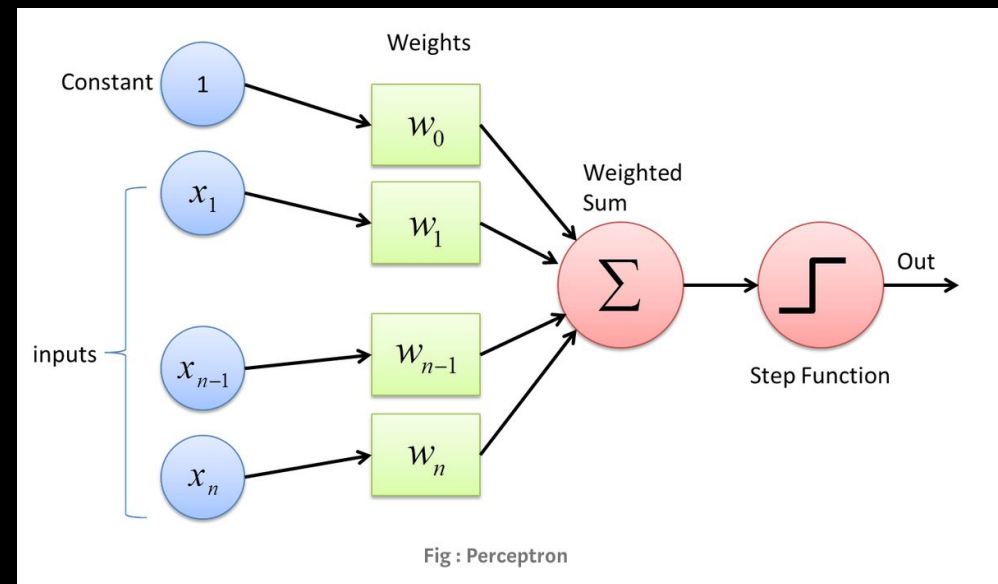

- (Highly) parametric non-linear function approximators.
- Loosely inspired from brain : network of interconnected neurons that “activate” or not.
- Composition of stacked ‘*layers*’ of weights (parameters) producing *neurons*.



Let's start with a very simple one: Perceptron

- Like linear regression: $\sum_i w_i x_i$
- But with an additional **activation function**

Here : “step” function, threshold to obtain binary outputs.




‘Vanilla’ Neural Networks : Multi Layer Perceptron (MLP)

- Also called ‘**fully-connected**’ **feed-forward** neural networks.
- Stack ‘perceptron’ on several *layers*.
- Each *hidden layer* has a *latent size* (output size)

Glossary:


- Input nodes
- Output nodes
- Connections / weights
- Activation function
- Hidden layers




[Playground.tensorflow.org](https://playground.tensorflow.org)

(Some) Activation Functions:


- Sigmoid:



- Tanh



- Rectified Linear Unit (ReLU) / LeakyReLU:



- Softmax : all neurons in $[0,1]$ and sum to 1 (good for probability distribution).

Training : Finding the “best” weights (= parameters)

Remember that ‘cost’ function (criterion) ?

=> Goal is to minimize that

How ? Using gradient descent and backpropagation of the gradients !

(So you need -almost- differentiable functions)

Learning algorithm

- “Forward pass”: Feed an example (or more) to the network
- Compute error with regard to your criterion (compared to expected targets)
- Compute gradients (wait for it)
- Update weights
- Repeat until stopping criterion

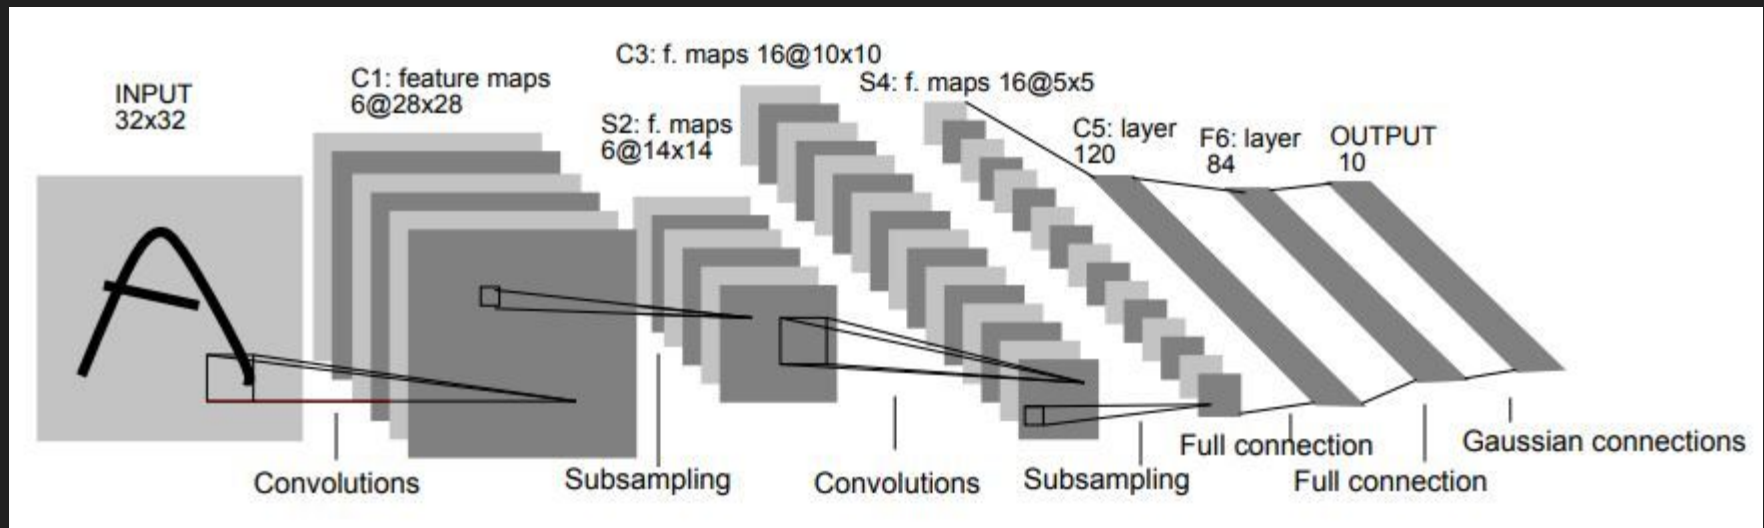
Backpropagation ?

- Goal of gradient : update the weights in the “right” direction
 - Which weight is responsible for the error and on what ‘amount’ ?
- Backpropagation relies on the chain rules to compute the gradients of a layer’s weight using the delta’s of next layer.

Different Losses



- Regression :
 - Mean-square error (MSE)
 - (Smooth) L1-Loss
- Classification :
 - Negative Log-Likelihood
 - Cross-Entropy
- And many others defined for specific problems...

Convolutional Neural Networks





Convolutional Neural Networks

- Use local connections instead of fully-connected
- Shared weights
- Each layer can train several convolution filters : channels




- Filters extend the full depth of the input volume
- Convolve the filter with the image: Slide over the image spatially, computing dot products.
- Stack the activation map produced by each filter : new “image”



Credit Fei-Fei Li's course


Pooling

- Number of features extracted can be very large (increasing number of channels / filters)
- To reduce the size of the activation maps, one can :
 - Use strides
 - Downsample using Max or AveragePooling - applied on each channel independently




Pooled feature (max
or mean operator)


Convolutional features



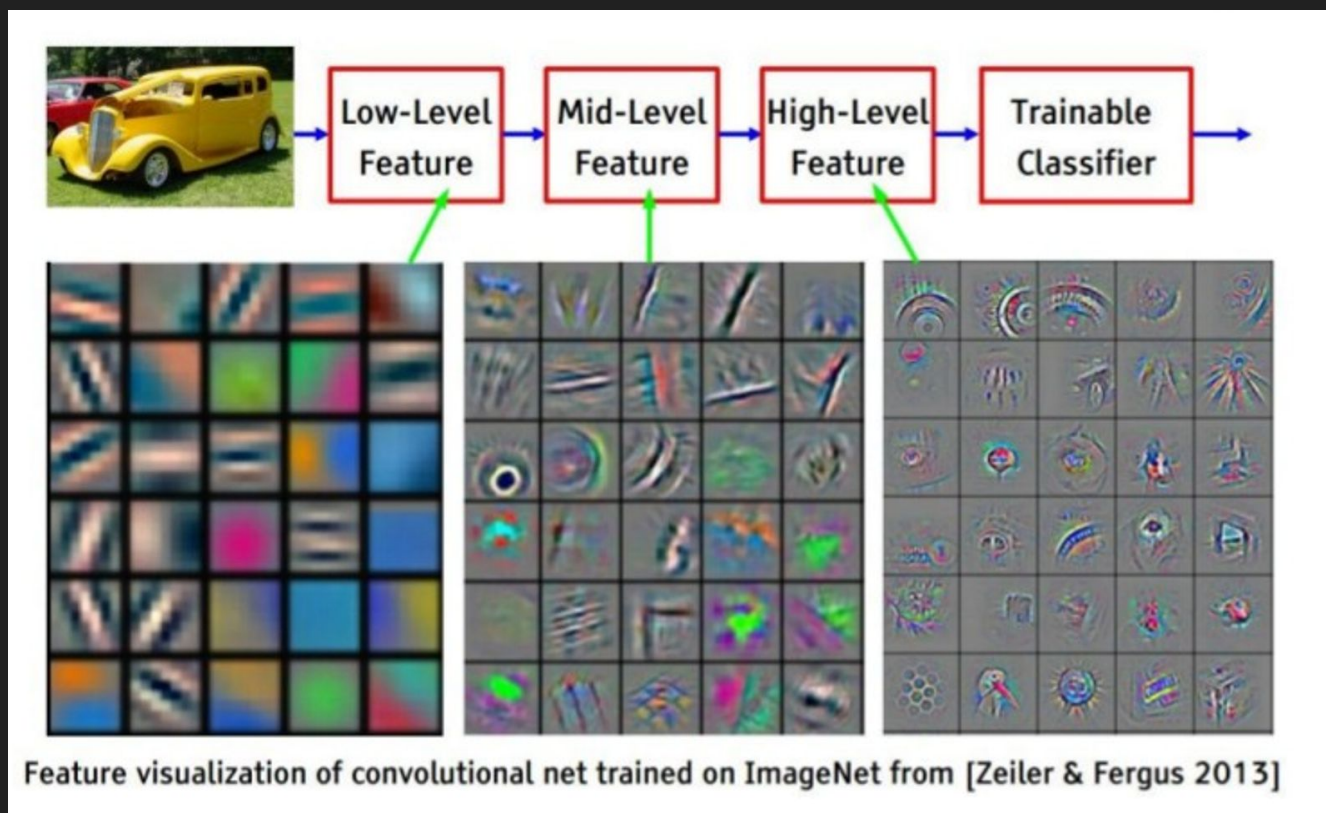
Credit Fei-Fei Li's course

Convolutional Neural Networks



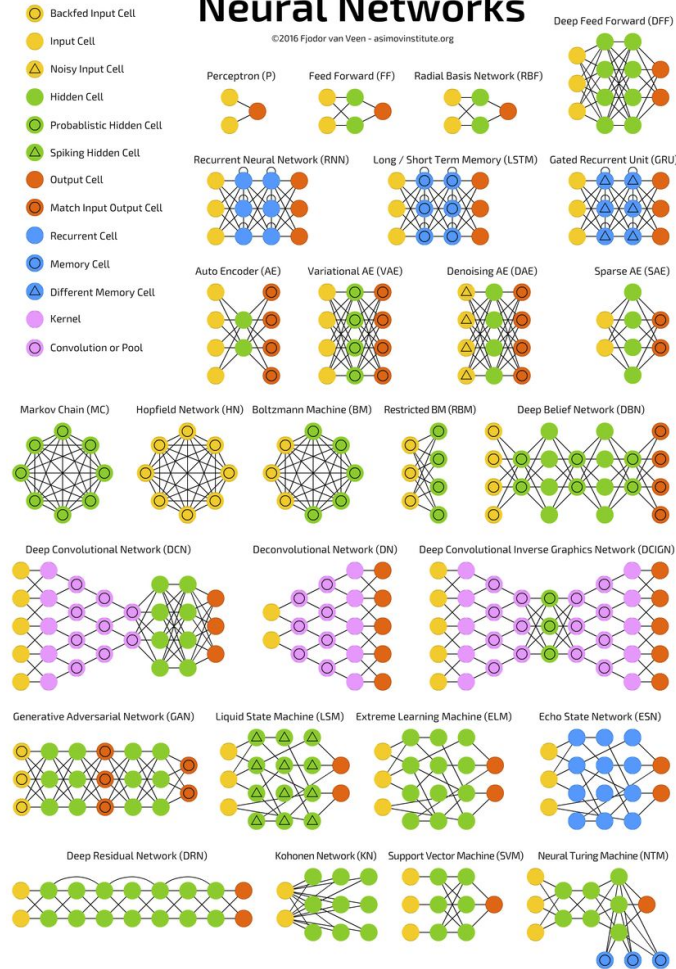


“Inside” deep learning : Representation Learning



A mostly complete chart of Neural Networks

©2016 Fjodor van Veen - asimovinstitute.org



Let's code !